Relevant models

AI models worth knowing about if you build with AI, and why each one matters.

Last updated: 13 Nov 2026 (Prague) · **How the ranking works:** what matters to a working AI developer: capability, whether you can call it through an API today, price for the performance, and how new it is. Vendor benchmark claims are the vendors' own.

What changed lately

- **Gemini 4 Hydrogen (12 Nov):** Google's new fast model beats Argon on the AA Intelligence Index at $1.50/$7.50 per M tokens and is available on all Gemini API tiers.
- **Inkling 2 Small weights out (12 Nov):** Thinking Machines released the Apache 2.0 weights on Hugging Face, plus API access through NVIDIA, Fireworks, Together, Baseten, Nebius, Modal and Cloudflare.
- **Gemini Decision 1 (11 Nov):** Google launched the Gemini Decisions API, joining Jev and OpenAI in decision models: up to 1,000 yes/no, choice or rubric questions per multimodal input in one pass, about 200 ms, $0.10 per M input tokens with free output.
- **Claude Fable 5.2 / Mythos 5.2 (9 Nov):** Anthropic's top model got faster (up to 30% vs 5.1) and cheaper ($8/$40). A swarm of about 1,000 Fable agents formalized Fermat's Last Theorem in Lean in 30 hours.
- **Muse Ember (9 Nov):** Meta's small agent workhorse costs $0.10/$0.50, has a 1M-token context, and Meta claims it beats Haiku 5.5 and GPT-6 Luna on intelligence per dollar.

Ranked

#	Model	Vendor	Released	Best for	Why you should care	Source
1	**Claude Fable 5.2**	Anthropic	9 Nov 2026	Mathematicians, the hardest open-ended technical work	It's the model for heavy reasoning and math: 1,000 Fable agents formalized Fermat's Last Theorem in Lean in 30 h (Opus 5.5 managed 26% in 18 h), and it solved two open black-hole physics problems. Same model as Mythos 5.2 but with cyber/bio safeguards. Up to 30% faster than 5.1, $8/$40 per M tokens (fast mode $16/$80), `claude-fable-5-2`, also on AWS, Google Cloud and Azure.	anthropic.com
2	**Claude Opus 5.5**	Anthropic	22 Sep 2026	Everyday coding and agent work	Anthropic's recommended daily driver: about Fable 5.1 level at $4/$20 with $0.20 cache reads, and 40% cheaper to run than Opus 5.	anthropic.com
3	**GPT-6.1 Sol**	OpenAI	29 Sep 2026	Coding, computer use, pro work at scale	Nearly matches Astra for a fifth of the price: $2/$10, $0.10 cached input.	openai.com
4	**GPT-6 Astra**	OpenAI	3 Sep 2026	Maximum capability, computer use, science	OpenAI's most capable model: 98% on FrontierMath Tier 4 and helped prove new prime-gap bounds. It costs $10/$50 per M tokens.	openai.com
5	**Gemini 4 Hydrogen**	Google	12 Nov 2026	Fast everyday coding and agent execution	Google's fast workhorse that edges out Argon on the AA Intelligence Index (51 vs 50 at default, 54 at xhigh) at a flat $1.50/$7.50 across its 1M context, with up to 1M output and four thinking levels. `gemini-4-hydrogen` on the Gemini API and AI Studio.	blog.google
6	**Gemini 4 Argon**	Google	13 Oct 2026	Long-horizon engineering, multimodal input	Google's GA frontier model: 1M unified input/output window with long decode. Intro price $2/$10 through 30 Nov, then $4/$20.	aistudio.google.com
7	**GPT-6.1 Terra**	OpenAI	Oct 2026	High-volume computer use and office automation	Fast and cheap at $1/$5; on par with Astra on Agents' Last Exam at a tenth of Astra's per-token price.	openai.com

#	Model	Vendor	Released	Best for	Why you should care	Source
8	**Kimi K3.1 Pro**	Moonshot AI	2 Nov 2026 (weights)	Open-weight frontier agentic coding	The biggest open-weight model here: 2.8T MoE / 104B active, 1M context, text+image. API as `kimi-k3.1-pro`; custom license.	huggingface.co
9	**DeepSeek-V4.1-Flash-1017**	DeepSeek	17 Oct 2026	Cheap strong open model, self-hosting	MIT-licensed 763B weights with a 1M context; it posts strong DeepSWE/Terminal-Bench scores on DeepSeek's own tables.	huggingface.co
10	**Jev 1.14**	TypeSafe AI	14 Oct 2026	Classifying, routing and scoring stuff in under a second	It's not a chatbot but a fast classifier: send it a state plus typed questions (Choice, Score, yes/no "Noul", Box) and get calibrated probabilities back in 70–500 ms instead of generated text. It never returns a type error, costs $0.042 per M input tokens, and output is free.	docs.typesafe.ai
11	**Gemini Decision 1**	Google	11 Nov 2026	Fast yes/no, choice and rubric decisions over any media	Google's answer to Jev: send one input (text, images, audio, video or PDF, up to 256K tokens) plus up to 1,000 typed questions and get calibrated probabilities back in about 150–250 ms. $0.10 per M input tokens, output free, via the Gemini Decisions API.	blog.google AI Studio docs
12	**Muse Ember**	Meta	9 Nov 2026	Subagents: tool calls, extraction, routing	Very cheap at $0.10/$0.50, with a 1M context and up to 512K output; Meta says it gives more intelligence per dollar than Haiku 5.5 or GPT-6 Luna. `muse-ember-1.0` on the Meta Model API.	research.meta.ai
13	**Inkling 2 Small**	Thinking Machines Lab	5 Nov 2026	Open-weight fine-tuning base, agentic coding	594B/17B MoE with image input and a 1M context, built to be fine-tuned. Apache 2.0 weights are now on Hugging Face (NVFP4 needs about 500GB of VRAM), with API access at NVIDIA, Fireworks, Together and others.	thinkingmachines.ai huggingface.co
14	**Claude Haiku 5.5**	Anthropic	7 Oct 2026	Fast, cheap subagents, support, browser use	$0.10/$0.50 for prompts up to 100K tokens, about 75% cheaper to run than Haiku 4.5, now with adjustable effort.	anthropic.com
15	**Mistral Large 4 ("Le Chonk")**	Mistral AI	Oct 2026 (weights)	Sovereign/self-hosted coding and cyber defense	European-trained ~1T open-weight model, $0.70/$2.10 via API. Free to use commercially unless you sell it as a model service.	mistral.ai
16	**Gemini 4 Carbon**	Google	28 Oct 2026	Cyber defense, very long agentic runs	Google's strongest model (#1 on AA Intelligence Index v5 at 59), but for now only cyber defenders in the Fairwind Program can use it.	blog.google
17	**GPT-6 Cyber**	OpenAI	5 Nov 2026	Vulnerability research, exploit validation	Found 81 high-severity Firefox bugs and 948 Linux exploits; access is limited to approved Daybreak Red customers.	openai.com
18	**Muse Wonder**	Meta	12 Oct 2026	Personal/professional agents with memory	Meta's frontier model, built around Active Memory and Active Focus, on the Meta Model API and Muse Code.	research.meta.ai
19	**Grok 4.7**	xAI (SpaceXSI)	21 Sep 2026	Coding, agents, knowledge work	xAI's current API flagship: 500K context at $2/$6. Grok 4.8 isn't on the API yet.	docs.x.ai
20	**Qwen4-Max-Preview**	Alibaba (Qwen)	11 Nov 2026 (preview)	Alibaba's flagship reasoning, long video	First Qwen 4 model: 1M context, up to 393K output, text/image/video input. Token Plan only for now.	docs.qwencloud.com
21	**GPT-6 Luna**	OpenAI	22 Sep 2026	Cheapest OpenAI workhorse	$0.10/$0.50; OpenAI pitches it for fast, efficient everyday work at scale.	openai.com

Prices are USD per million input/output tokens at standard API rates unless noted.

Open weights worth running

Downloadable weights; the note gives the license.

- **Kimi K3.1 Pro**: open now, custom license
- **DeepSeek-V4.1-Flash-1017**: open now, MIT
- **Inkling 2 Small**: open now, Apache 2.0
- **Mistral Large 4 ("Le Chonk")**: open now, modified MIT; Model-as-a-Service use needs a commercial license

Fast/cheap specialists

- **Jev 1.14**: classify and route in under a second
- **Gemini Decision 1**: multimodal decisions in ~200 ms
- **Muse Ember**
- **Claude Haiku 5.5**
- **GPT-6 Luna**

Refreshed on weekday mornings (Prague). Models are re-ranked by developer relevance (capability, API availability, price/performance, novelty), and every entry links a primary source checked at refresh time.

Source repository: https://github.com/janlycka/relevant-models