Relevant models

AI models worth knowing about if you build with AI, and why each one matters.

Last updated: 12 Nov 2026 (Prague) · **How the ranking works:** what matters to a working AI developer: capability, whether you can call it through an API today, price for the performance, and how new it is. Vendor benchmark claims are the vendors' own.

What changed lately

- **Gemini Decision 1 (11 Nov):** Google launched the Gemini Decisions API, joining Jev and OpenAI in decision models: up to 1,000 yes/no, choice or rubric questions per multimodal input in one pass, about 200 ms, $0.10 per M input tokens with free output.
- **Claude Fable 5.2 / Mythos 5.2 (9 Nov):** Anthropic's top model got faster (up to 30% vs 5.1) and cheaper ($8/$40). A swarm of about 1,000 Fable agents formalized Fermat's Last Theorem in Lean in 30 hours.
- **Muse Ember (9 Nov):** Meta's small agent workhorse costs $0.10/$0.50, has a 1M-token context, and Meta claims it beats Haiku 5.5 and GPT-6 Luna on intelligence per dollar.
- **Inkling 2 Small (5 Nov):** Thinking Machines' 594B-total/17B-active MoE. It runs on Tinker now; the Apache 2.0 weights are promised after safety testing and were not on Hugging Face as of 11 Nov.
- **Gated cyber models:** GPT-6 Cyber (5 Nov) and Gemini 4 Carbon (28 Oct) are both out, but only to vetted defenders for now.

Ranked

#	Model	Vendor	Released	Best for	Why you should care	Source
1	**Claude Fable 5.2**	Anthropic	9 Nov 2026	Mathematicians, the hardest open-ended technical work	It's the model for heavy reasoning and math: 1,000 Fable agents formalized Fermat's Last Theorem in Lean in 30 h (Opus 5.5 managed 26% in 18 h), and it solved two open black-hole physics problems. Same model as Mythos 5.2 but with cyber/bio safeguards. Up to 30% faster than 5.1, $8/$40 per M tokens (fast mode $16/$80), `claude-fable-5-2`, also on AWS, Google Cloud and Azure.	anthropic.com
2	**Claude Opus 5.5**	Anthropic	22 Sep 2026	Everyday coding and agent work	Anthropic's recommended daily driver: about Fable 5.1 level at $4/$20 with $0.20 cache reads, and 40% cheaper to run than Opus 5.	anthropic.com
3	**GPT-6.1 Sol**	OpenAI	29 Sep 2026	Coding, computer use, pro work at scale	Nearly matches Astra for a fifth of the price: $2/$10, $0.10 cached input.	openai.com
4	**GPT-6 Astra**	OpenAI	3 Sep 2026	Maximum capability, computer use, science	OpenAI's most capable model: 98% on FrontierMath Tier 4 and helped prove new prime-gap bounds. It costs $10/$50 per M tokens.	openai.com
5	**Gemini 4 Argon**	Google	13 Oct 2026	Long-horizon engineering, multimodal input	Google's GA frontier model: 1M unified input/output window with long decode. Intro price $2/$10 through 30 Nov, then $4/$20.	aistudio.google.com
6	**GPT-6.1 Terra**	OpenAI	Oct 2026	High-volume computer use and office automation	Fast and cheap at $1/$5; on par with Astra on Agents' Last Exam at a tenth of Astra's per-token price.	openai.com
7	**Kimi K3.1 Pro**	Moonshot AI	2 Nov 2026 (weights)	Open-weight frontier agentic coding	The biggest open-weight model here: 2.8T MoE / 104B active, 1M context, text+image. API as `kimi-k3.1-pro`; custom license.	huggingface.co
8	**DeepSeek-V4.1-Flash-1017**	DeepSeek	17 Oct 2026	Cheap strong open model, self-hosting	MIT-licensed 763B weights with a 1M context; it posts strong DeepSWE/Terminal-Bench scores on DeepSeek's own tables.	huggingface.co

#	Model	Vendor	Released	Best for	Why you should care	Source
9	**Jev 1.14**	TypeSafe AI	14 Oct 2026	Classifying, routing and scoring stuff in under a second	It's not a chatbot but a fast classifier: send it a state plus typed questions (Choice, Score, yes/no "Noul", Box) and get calibrated probabilities back in 70–500 ms instead of generated text. It never returns a type error, costs $0.042 per M input tokens, and output is free.	docs.typesafe.ai
10	**Gemini Decision 1**	Google	11 Nov 2026	Fast yes/no, choice and rubric decisions over any media	Google's answer to Jev: send one input (text, images, audio, video or PDF, up to 256K tokens) plus up to 1,000 typed questions and get calibrated probabilities back in about 150–250 ms. $0.10 per M input tokens, output free, via the Gemini Decisions API.	blog.google AI Studio docs
11	**Muse Ember**	Meta	9 Nov 2026	Subagents: tool calls, extraction, routing	Very cheap at $0.10/$0.50, with a 1M context and up to 512K output; Meta says it gives more intelligence per dollar than Haiku 5.5 or GPT-6 Luna. `muse-ember-1.0` on the Meta Model API.	research.meta.ai
12	**Claude Haiku 5.5**	Anthropic	7 Oct 2026	Fast, cheap subagents, support, browser use	$0.10/$0.50 for prompts up to 100K tokens, about 75% cheaper to run than Haiku 4.5, now with adjustable effort.	anthropic.com
13	**Mistral Large 4 ("Le Chonk")**	Mistral AI	Oct 2026 (weights)	Sovereign/self-hosted coding and cyber defense	European-trained ~1T open-weight model, $0.70/$2.10 via API. Free to use commercially unless you sell it as a model service.	mistral.ai
14	**Gemini 4 Carbon**	Google	28 Oct 2026	Cyber defense, very long agentic runs	Google's strongest model (#1 on AA Intelligence Index v5 at 59), but for now only cyber defenders in the Fairwind Program can use it.	blog.google
15	**GPT-6 Cyber**	OpenAI	5 Nov 2026	Vulnerability research, exploit validation	Found 81 high-severity Firefox bugs and 948 Linux exploits; access is limited to approved Daybreak Red customers.	openai.com
16	**Muse Wonder**	Meta	12 Oct 2026	Personal/professional agents with memory	Meta's frontier model, built around Active Memory and Active Focus, on the Meta Model API and Muse Code.	research.meta.ai
17	**Inkling 2 Small**	Thinking Machines Lab	5 Nov 2026	Fine-tuning base, agentic coding	594B/17B MoE built to be fine-tuned on Tinker. Apache 2.0 weights are promised but not out yet.	thinkingmachines.ai
18	**Grok 4.7**	xAI (SpaceXSI)	21 Sep 2026	Coding, agents, knowledge work	xAI's current API flagship: 500K context at $2/$6. Grok 4.8 isn't on the API yet.	docs.x.ai
19	**Qwen4-Max-Preview**	Alibaba (Qwen)	11 Nov 2026 (preview)	Alibaba's flagship reasoning, long video	First Qwen 4 model: 1M context, up to 393K output, text/image/video input. Token Plan only for now.	docs.qwencloud.com
20	**Gemini Omni 1.2 Flash**	Google	5 Nov 2026	Video generation	GA video model: 15-second clips, multi-turn editing, $0.07/s at 720p (30% cheaper than 1.1).	blog.google
21	**GPT-6 Luna**	OpenAI	22 Sep 2026	Cheapest OpenAI workhorse	$0.10/$0.50; OpenAI pitches it for fast, efficient everyday work at scale.	openai.com

Prices are USD per million input/output tokens at standard API rates unless noted.

Open weights worth running

Downloadable weights; the note gives the license.

- **Kimi K3.1 Pro**: open now, custom license
- **DeepSeek-V4.1-Flash-1017**: open now, MIT
- **Mistral Large 4 ("Le Chonk")**: open now, modified MIT; Model-as-a-Service use needs a commercial license

- **Inkling 2 Small**: pending: Apache 2.0 weights promised after safety testing

Fast/cheap specialists

- **Jev 1.14**: classify and route in under a second
- **Gemini Decision 1**: multimodal decisions in ~200 ms
- **Muse Ember**
- **Claude Haiku 5.5**
- **GPT-6 Luna**

Refreshed on weekday mornings (Prague). Models are re-ranked by developer relevance (capability, API availability, price/performance, novelty), and every entry links a primary source checked at refresh time.

Source repository: https://github.com/janlycka/relevant-models